

ml-assignment-4016

April 12, 2023

1 Simple Linear Regression

```
[134]: import matplotlib.pyplot as plt  
import pandas as pd  
import numpy as np
```

```
[135]: df = pd.read_csv('/content/placement.csv')  
df.head()
```

```
[135]:
```

	cgpa	package
0	6.89	3.26
1	5.12	1.98
2	7.82	3.25
3	7.42	3.67
4	6.94	3.57

```
[136]: plt.scatter(df['cgpa'],df['package'])  
plt.xlabel('CGPA')  
plt.ylabel('Package(in lpa)')
```

```
[136]: Text(0, 0.5, 'Package(in lpa)')
```



```
[137]: df.describe()
```

```
[137]:
```

	cgpa	package
count	200.000000	200.000000
mean	6.990500	2.996050
std	1.069409	0.691644
min	4.260000	1.370000
25%	6.190000	2.487500
50%	6.965000	2.995000
75%	7.737500	3.492500
max	9.580000	4.620000

```
[138]: X = df.iloc[:,0:1]
y = df.iloc[:, -1]
```

```
[139]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
↪2,random_state=2)
```

```
[140]: from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
```

```
[140]: LinearRegression()
```

```
[141]: model.coef_, model.intercept_
```

```
[141]: (array([0.55795197]), -0.8961119222429144)
```


```
[142]: df.plot.scatter('cgpa', 'package', color='blue')
plt.plot(
    np.array(df['cgpa']),
    model.coef_[0] * np.array(df['cgpa']) + model.intercept_,
    color='red'
)
plt.show()
```



```
[143]: from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.  
↪2,random_state=2)
```

```
[144]: X_test.plot.density()  
plt.show()
```



```
[145]: model.predict(X_test.iloc[0].values.reshape(1,1))
```

/usr/local/lib/python3.9/dist-packages/sklearn/base.py:439: UserWarning:

X does not have valid feature names, but LinearRegression was fitted with
feature names

```
[145]: array([3.89111601])
```

```
[147]: m = model.coef_  
b = model.intercept_  
print(m,b)
```

```
[0.55795197] -0.8961119222429144
```

2 Linear Regression with Multiple Variable

```
[105]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
sns.set_theme(style='darkgrid')
```

```
[92]: df = pd.read_csv('/content/houseprices2.csv')
df.head()
```

```
[92]:
```

	area	bedrooms	age	price
0	2600	3.0	20	550000
1	3000	4.0	15	565000
2	3200	NaN	18	610000
3	3600	3.0	30	595000
4	4000	5.0	8	760000

```
[93]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   area        6 non-null      int64
1   bedrooms    5 non-null      float64
2   age         6 non-null      int64
3   price       6 non-null      int64
dtypes: float64(1), int64(3)
memory usage: 320.0 bytes
```

```
[94]: df.describe()
```

```
[94]:
```

	area	bedrooms	age	price
count	6.000000	5.00000	6.000000	6.000000
mean	3416.666667	4.20000	16.500000	648333.333333
std	587.934237	1.30384	8.288546	109117.673484
min	2600.000000	3.00000	8.000000	550000.000000
25%	3050.000000	3.00000	9.750000	572500.000000
50%	3400.000000	4.00000	16.500000	602500.000000
75%	3900.000000	5.00000	19.500000	722500.000000
max	4100.000000	6.00000	30.000000	810000.000000

```
[106]: fig = px.scatter_3d(df, x='area', y='bedrooms', z='age')
```


```
fig.show()
```


```
[95]: sns.heatmap(df.corr(), annot=True)
```


```
[95]: <Axes: >
```



```
[96]: for x in df.columns[:-1]:  
       df.plot.scatter(x, 'price', color='red')  
       plt.show()
```










```
[97]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()


lr.fit(X_train,y_train)
```

```
[97]: LinearRegression()
```

```
[98]: for x in df.columns:
    df[x].plot.density()
    plt.show()
```







```
[101]: df['bedrooms'] = df['bedrooms'].fillna(
        df['bedrooms'].median()
    )
    df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   area        6 non-null      int64
 1   bedrooms    6 non-null      float64
 2   age         6 non-null      int64
 3   price       6 non-null      int64
dtypes: float64(1), int64(3)
memory usage: 320.0 bytes
```

```
[102]: X_train = df.drop('price', axis=1)
        y_train = df['price']
```

```
[103]: from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
```

```
[103]: LinearRegression()
```

```
[104]: model.coef_, model.intercept_
```

```
[104]: (array([ 112.06244194, 23388.88007794, -3231.71790863]), 221323.00186540396)
```

3 Polynomial Regression

```
[132]: import math
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
sns.set_theme(style='darkgrid')
```

```
[148]: df = pd.read_csv('/content/Position_Salaries.csv')
df
```

```
[148]:
```

	Position	Level	Salary
0	Business Analyst	1	45000
1	Junior Consultant	2	50000
2	Senior Consultant	3	60000
3	Manager	4	80000
4	Country Manager	5	110000
5	Region Manager	6	150000
6	Partner	7	200000
7	Senior Partner	8	300000
8	C-level	9	500000
9	CEO	10	1000000

```
[149]: df.describe()
```

```
[149]:
```

	Level	Salary
count	10.00000	10.000000
mean	5.50000	249500.000000
std	3.02765	299373.883668
min	1.00000	45000.000000
25%	3.25000	65000.000000
50%	5.50000	130000.000000
75%	7.75000	275000.000000
max	10.00000	1000000.000000


```
[150]: df = df.drop('Position', axis=1)
df.head()
```

```
[150]:   Level  Salary
0      1   45000
1      2   50000
2      3   60000
3      4   80000
4      5  110000
```

```
[151]: sns.heatmap(df.corr(), annot=True)
plt.show()
```



```
[152]: df.plot.scatter('Level', 'Salary', color='red')
plt.show()
```



```
[153]: X_train = df.drop('Salary', axis=1)
       y_train = df['Salary']
```

- 1 Degree

```
[154]: from sklearn.linear_model import LinearRegression
       model = LinearRegression()
       model.fit(X_train, y_train)
```


```
[154]: LinearRegression()
```

```
[155]: model.coef_, model.intercept_
```

```
[155]: (array([80878.78787879]), -195333.33333333337)
```

```
[156]: df.plot.scatter('Level', 'Salary', color='blue')
       plt.plot(
           np.array(df['Level']),
           model.coef_[0] * np.array(df['Level']) + model.intercept_,
           color='red')
```

```
)  
plt.show()
```



- 2 Degree

```
[157]: from sklearn.preprocessing import PolynomialFeatures  
quadratic = PolynomialFeatures(degree=2)  
X_deg2 = quadratic.fit_transform(X_train)  
X_deg2
```

```
[157]: array([[ 1.,  1.,  1.],  
[ 1.,  2.,  4.],  
[ 1.,  3.,  9.],  
[ 1.,  4., 16.],  
[ 1.,  5., 25.],  
[ 1.,  6., 36.],  
[ 1.,  7., 49.],  
[ 1.,  8., 64.],  
[ 1.,  9., 81.],  
[ 1., 10., 100.]])
```




```
[158]: model_deg2 = LinearRegression()  
model_deg2.fit(X_deg2, y_train)
```

```
[158]: LinearRegression()
```

```
[159]: model_deg2.coef_, model_deg2.intercept_
```

```
[159]: (array([ 0.          , -132871.21212121,  19431.81818182]),  
       232166.6666666664)
```

```
[160]: df.plot.scatter('Level', 'Salary', color='blue')  
plt.plot(  
    np.array(df['Level']),  
    model_deg2.coef_[1] * np.array(df['Level'])  
    + model_deg2.coef_[2] * np.array(df['Level']) ** 2  
    + model_deg2.intercept_,  
    color='red'  
)  
plt.show()
```



- 3 Degree

```
[161]: cubic = PolynomialFeatures(degree=3)
X_deg3 = cubic.fit_transform(X_train)
X_deg3
```

```
[161]: array([[ 1.,   1.,   1.,   1.],
 [ 1.,   2.,   4.,   8.],
 [ 1.,   3.,   9.,  27.],
 [ 1.,   4.,  16.,  64.],
 [ 1.,   5.,  25., 125.],
 [ 1.,   6.,  36., 216.],
 [ 1.,   7.,  49., 343.],
 [ 1.,   8.,  64., 512.],
 [ 1.,   9.,  81., 729.],
 [ 1.,  10., 100., 1000.]])
```


```
[162]: model_deg3 = LinearRegression()
model_deg3.fit(X_deg3, y_train)
```

```
[162]: LinearRegression()
```

```
[163]: model_deg3.coef_, model_deg3.intercept_
```

```
[163]: (array([ 0.          , 180664.33566432, -48548.95104895,  4120.04662005]),
 -121333.33333330264)
```

```
[164]: df.plot.scatter('Level', 'Salary', color='blue')
plt.plot(
    np.array(df['Level']),
    model_deg3.coef_[1] * np.array(df['Level'])
    + model_deg3.coef_[2] * np.array(df['Level']) ** 2
    + model_deg3.coef_[3] * np.array(df['Level']) ** 3
    + model_deg3.intercept_,
    color='red'
)
plt.show()
```



- 4 Degree

```
[165]: quartic = PolynomialFeatures(degree=4)
X_deg4 = quartic.fit_transform(X_train)
X_deg4
```

```
[165]: array([[1.000e+00, 1.000e+00, 1.000e+00, 1.000e+00, 1.000e+00],
 [1.000e+00, 2.000e+00, 4.000e+00, 8.000e+00, 1.600e+01],
 [1.000e+00, 3.000e+00, 9.000e+00, 2.700e+01, 8.100e+01],
 [1.000e+00, 4.000e+00, 1.600e+01, 6.400e+01, 2.560e+02],
 [1.000e+00, 5.000e+00, 2.500e+01, 1.250e+02, 6.250e+02],
 [1.000e+00, 6.000e+00, 3.600e+01, 2.160e+02, 1.296e+03],
 [1.000e+00, 7.000e+00, 4.900e+01, 3.430e+02, 2.401e+03],
 [1.000e+00, 8.000e+00, 6.400e+01, 5.120e+02, 4.096e+03],
 [1.000e+00, 9.000e+00, 8.100e+01, 7.290e+02, 6.561e+03],
 [1.000e+00, 1.000e+01, 1.000e+02, 1.000e+03, 1.000e+04]])
```


```
[166]: model_deg4 = LinearRegression()
model_deg4.fit(X_deg4, y_train)
```

```
[166]: LinearRegression()
```

```
[167]: model_deg4.coef_, model_deg4.intercept_
```

```
[167]: (array([      0.          , -211002.33100292,  94765.44289063,  
          -15463.28671331,    890.15151515]),  
       184166.66666719783)
```

```
[168]: df.plot.scatter('Level', 'Salary', color='blue')  
plt.plot(  
    np.array(df['Level']),  
    model_deg4.coef_[1] * np.array(df['Level'])  
        + model_deg4.coef_[2] * np.array(df['Level']) ** 2  
        + model_deg4.coef_[3] * np.array(df['Level']) ** 3  
        + model_deg4.coef_[4] * np.array(df['Level']) ** 4  
        + model_deg4.intercept_,  
    color='red'  
)  
plt.show()
```



4 Linear Regression with Gradient Descent

```
[107]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
sns.set_theme(style='darkgrid')
```

```
[108]: df = pd.read_csv('https://github.com/codebasics/py/raw/master/ML/1_linear_reg/
↳homeprices.csv')
df.head()
```


```
[108]:   area  price
0  2600  550000
1  3000  565000
2  3200  610000
3  3600  680000
4  4000  725000
```

```
[109]: df.describe()
```

```
[109]:
```

	area	price
count	5.000000	5.000000
mean	3280.000000	626000.000000
std	540.370243	74949.983322
min	2600.000000	550000.000000
25%	3000.000000	565000.000000
50%	3200.000000	610000.000000
75%	3600.000000	680000.000000
max	4000.000000	725000.000000

```
[110]: df.plot.scatter('area', 'price', color='red')
plt.show()
```



```
[111]: X_train = df.drop('price', axis=1)
       y_train = df.drop('area', axis=1)
```

```
[112]: from sklearn import preprocessing
       scaler = preprocessing.StandardScaler()
       X_train_scaled = scaler.fit_transform(X_train)
       y_train_scaled = scaler.fit_transform(y_train)
```

```
[113]: from sklearn.metrics import mean_squared_error

def gradient_descent(X, y, theta, lr=0.01, steps=100):
    N = len(y)
    cost_history = np.zeros(steps)
    theta_history = np.zeros((steps, 2))
    for i in range(steps):
        prediction = np.dot(X, theta)
        theta = theta - (1 / N) * lr * (X.T.dot((prediction - y)))
        theta_history[i, :] = theta.T
        cost_history[i] = mean_squared_error(prediction, y)


    return theta, cost_history, theta_history
```

```
[114]: lr = 0.1
it = 50
theta = np.random.randn(2, 1)

X_b = np.c_[
    X_train_scaled,
    np.ones((len(X_train_scaled), 1))
]

_, cost, theta = gradient_descent(X_b, y_train_scaled, theta, lr, it)
```

```
[115]: sns.lineplot(x=range(1, 51), y=cost)
plt.show()
```




```
[116]: np.array([1,2])
theta_df = pd.DataFrame(theta)
```


```
[117]: fig = plt.figure()
ax = plt.axes(projection='3d')

ax.plot_surface(theta_df[0], theta_df.drop(0, axis=1), np.array(cost).
    ↪reshape(-1, 1), cmap='viridis', edgecolor='none')
```


```
ax.set_title('Surface plot')  
plt.show()
```



```
[118]: sns.lineplot(x=theta_df[1], y=cost)  
plt.show()
```

```
[119]: plt.scatter(X_train_scaled, y_train_scaled, color='blue')
for t in theta:
    plt.plot(
        X_train_scaled,
        t[0] * X_train_scaled + t[1],
        color='red',
        alpha=0.1
    )
plt.show()
```




5 Logistic Regression

```
[120]: from sklearn.datasets import make_classification
import numpy as np
X, y = make_classification(n_samples=100, n_features=2,
    ↳ n_informative=1, n_redundant=0, n_classes=2, n_clusters_per_class=1,
    ↳ random_state=41, hypercube=False, class_sep=20)
```

```
[123]: import matplotlib.pyplot as plt
plt.figure(figsize=(10,6))
plt.scatter(X[:,0],X[:,1],c=y,cmap='winter',s=100)
```

```
[123]: <matplotlib.collections.PathCollection at 0x7f46626c2ee0>
```



```
[124]: from sklearn.linear_model import LogisticRegression
lor = LogisticRegression(penalty='none', solver='sag')
lor.fit(X,y)
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:1173:
FutureWarning:
```

```
`penalty='none'`has been deprecated in 1.2 and will be removed in 1.4. To keep
the past behaviour, set `penalty=None`.
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning:
```

```
The max_iter was reached which means the coef_ did not converge
```

```
[124]: LogisticRegression(penalty='none', solver='sag')
```

```
[125]: print(lor.coef_)
print(lor.intercept_)
```

```
[[4.79291636 0.20533276]]
[5.76753249]
```

```
[126]: m1 = -(lor.coef_[0][0]/lor.coef_[0][1])  
b1 = -(lor.intercept_/lor.coef_[0][1])
```

```
[127]: x_input = np.linspace(-3,3,100)  
y_input = m1*x_input + b1
```

```
[129]: def sigmoid(z):  
    return 1/(1 + np.exp(-z))
```

```
[128]: def gd(X,y):  
  
    X = np.insert(X,0,1,axis=1)  
    weights = np.ones(X.shape[1])  
    lr = 0.5  
  
    for i in range(5000):  
        y_hat = sigmoid(np.dot(X,weights))  
        weights = weights + lr*(np.dot((y-y_hat),X)/X.shape[0])  
  
    return weights[1:],weights[0]
```

```
[130]: coef_,intercept_ = gd(X,y)  
m = -(coef_[0]/coef_[1])  
b = -(intercept_/coef_[1])  
x_input1 = np.linspace(-3,3,100)  
y_input1 = m*x_input1 + b
```

```
[131]: plt.figure(figsize=(10,6))  
plt.plot(x_input,y_input,color='red',linewidth=3)  
plt.plot(x_input1,y_input1,color='black',linewidth=3)  
plt.scatter(X[:,0],X[:,1],c=y,cmap='winter',s=100)  
plt.ylim(-3,2)
```

```
[131]: (-3.0, 2.0)
```

